

HW2: Neurons/Learning Rate

Alex Shah

10/20/17

0.1 Code

```
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
import time

startTime = time.process_time()

from tensorflow.examples.tutorials.mnist import
    input_data
mnist = input_data.read_data_sets(".", one_hot=True)

trainimgs = mnist.train.images
trainlabels = mnist.train.labels
testimgs = mnist.test.images
testlabels = mnist.test.labels

ntrain = trainimgs.shape[0]
ntest = testimgs.shape[0]
dim = trainimgs.shape[1]
nclasses = trainlabels.shape[1]

## Change for HW2
n_hidden = 4096
learning_rate = 1.0
training_iters = 100000
batch_size = 100
##

n_input = 28 # MNIST data input (img shape: 28*28)
n_steps = 28 # timesteps
n_classes = 10 # MNIST total classes (0-9 digits)
display_step = 10

x = tf.placeholder(dtype="float", shape=[None, n_steps,
    n_input], name="x") # Current data input shape: (
    batch_size, n_steps, n_input) [100x28x28]
y = tf.placeholder(dtype="float", shape=[None, n_classes
    ], name="y")

#different way of writing out a dictionary, or variable
as a dictionary
weights = {
    'out': tf.Variable(tf.random_normal([n_hidden, n_classes
    ]))
}
biases = {
```

```

'out': tf.Variable(tf.random_normal([n_classes]))
}

lstm_cell = tf.contrib.rnn.BasicLSTMCell(n_hidden,
    forget_bias=1.0)

outputs, states = tf.nn.dynamic_rnn(lstm_cell, inputs=x,
    dtype=tf.float32)

output = tf.reshape(tf.split(outputs, 28, axis=1, num=
    None, name='split')[-1],[-1,n_hidden])
pred = tf.matmul(output, weights['out']) + biases['out']

cost = tf.reduce_mean(tf.nn.
    softmax_cross_entropy_with_logits(labels=y, logits=
    pred))
optimizer = tf.train.AdamOptimizer(learning_rate=
    learning_rate).minimize(cost)

#define accuracy for learning
correct_pred = tf.equal(tf.argmax(pred,1), tf.argmax(y,1)
)
accuracy = tf.reduce_mean(tf.cast(correct_pred, tf.
    float32))

init = tf.global_variables_initializer()

##Out of memory tweaks, didn't work for 4096
config = tf.ConfigProto()
#config.gpu_options.per_process_gpu_memory_fraction = 0.2
config.gpu_options.allow_growth = True
with tf.Session(config=config) as sess:
    sess.run(init)
    step = 1
    # Keep training until reach max iterations
    while step * batch_size < training_iters:

        # We will read a batch of 100 images [100 x 784]
        as batch_x
        # batch_y is a matrix of [100x10]
        batch_x, batch_y = mnist.train.next_batch(
            batch_size)

        # We consider each row of the image as one
        sequence
        # Reshape data to get 28 seq of 28 elements, so
        that, batch_x is [100x28x28]
        batch_x = batch_x.reshape((batch_size, n_steps,
            n_input))

```

```

# Run optimization op (backprop)
sess.run(optimizer, feed_dict={x: batch_x, y:
    batch_y})

if step % display_step == 0:
    # Calculate batch accuracy
    acc = sess.run(accuracy, feed_dict={x:
        batch_x, y: batch_y})
    # Calculate batch loss
    loss = sess.run(cost, feed_dict={x: batch_x,
        y: batch_y})
    step += 1

# Calculate accuracy for all mnist test images
test_data = mnist.test.images.reshape((-1, n_steps,
    n_input))
test_label = mnist.test.labels
print("Testing_Accuracy:", \
    sess.run(accuracy, feed_dict={x: test_data, y:
        test_label}))

sess.close()
currentTime = time.process_time()-startTime
print("Number_of_hidden_layers:", n_hidden)
print("Learning_Rate:", learning_rate)
print("Time_from_start:", currentTime)
print("=====")

```

0.2 Results

Table 1: Neurons and Learning Rate Test Accuracy

Neurons/Learning Rate	0.00001	0.0001	0.001	0.01	0.1	1.0
16	0.1083	0.3873	0.8545	0.9487	0.8425	0.3734
32	0.2175	0.6249	0.9438	0.9703	0.8836	0.0892
64	0.2156	0.8204	0.9619	0.9752	0.8127	0.0958
128	0.4487	0.903	0.9688	0.9776	0.4733	0.1009
256	0.6379	0.9276	0.9705	0.9786	0.098	0.1028
512	0.764	0.9398	0.9661	0.9768	0.4154	0.098
1024	0.8381	0.945	0.9711	0.3739	0.1032	0.1135
2048	0.8541	0.9435	0.9707	0.0892	0.0982	0.1135
4096	DNF	DNF	DNF	DNF	DNF	DNF

0.2.1 Explanation

What is this telling you about the training of an LSTM with respect to the number of neurons or learning rate or both? What can you learn from this?

We can see as the number of neurons increases the accuracy tends to rise. As another general case we can see that the best performance for learning rate is a "goldilocks" zone, varying with the neuron count, where the learning rate is not too small that the steps are inefficient, but not too large that it's leaping past finding it's mark.

What's also interesting is how neuron count affects a larger learning rate. A smaller number of neurons has better accuracy when paired with a larger learning rate, compared to larger number of neurons with the same larger learning rate. Conversely, as the number of neurons increases, we can see the best accuracy shift towards smaller learning rates.

A very large number of hidden units can take advantage of very tiny steps in the smallest learning rates. Whereas less hidden units requires a larger learning rate. My GPU could only run up to 2048 hidden units before running out of memory, and it took 122 seconds to complete each pass. The best accuracy it could accomplish with 2048 hidden units was 0.9707. Comparatively, 32 hidden units achieved a comparable accuracy of 0.9703 but only took 16 seconds to run. While increasing the hidden units was sometimes able exceed this value, the advantage to using a very high number of hidden units is lost to time and is computationally wasteful.